

The AES Process

1997: NIST publishes request for proposal


1998: 15 submissions. Five claimed attacks.

1999: NIST chooses 5 finalists

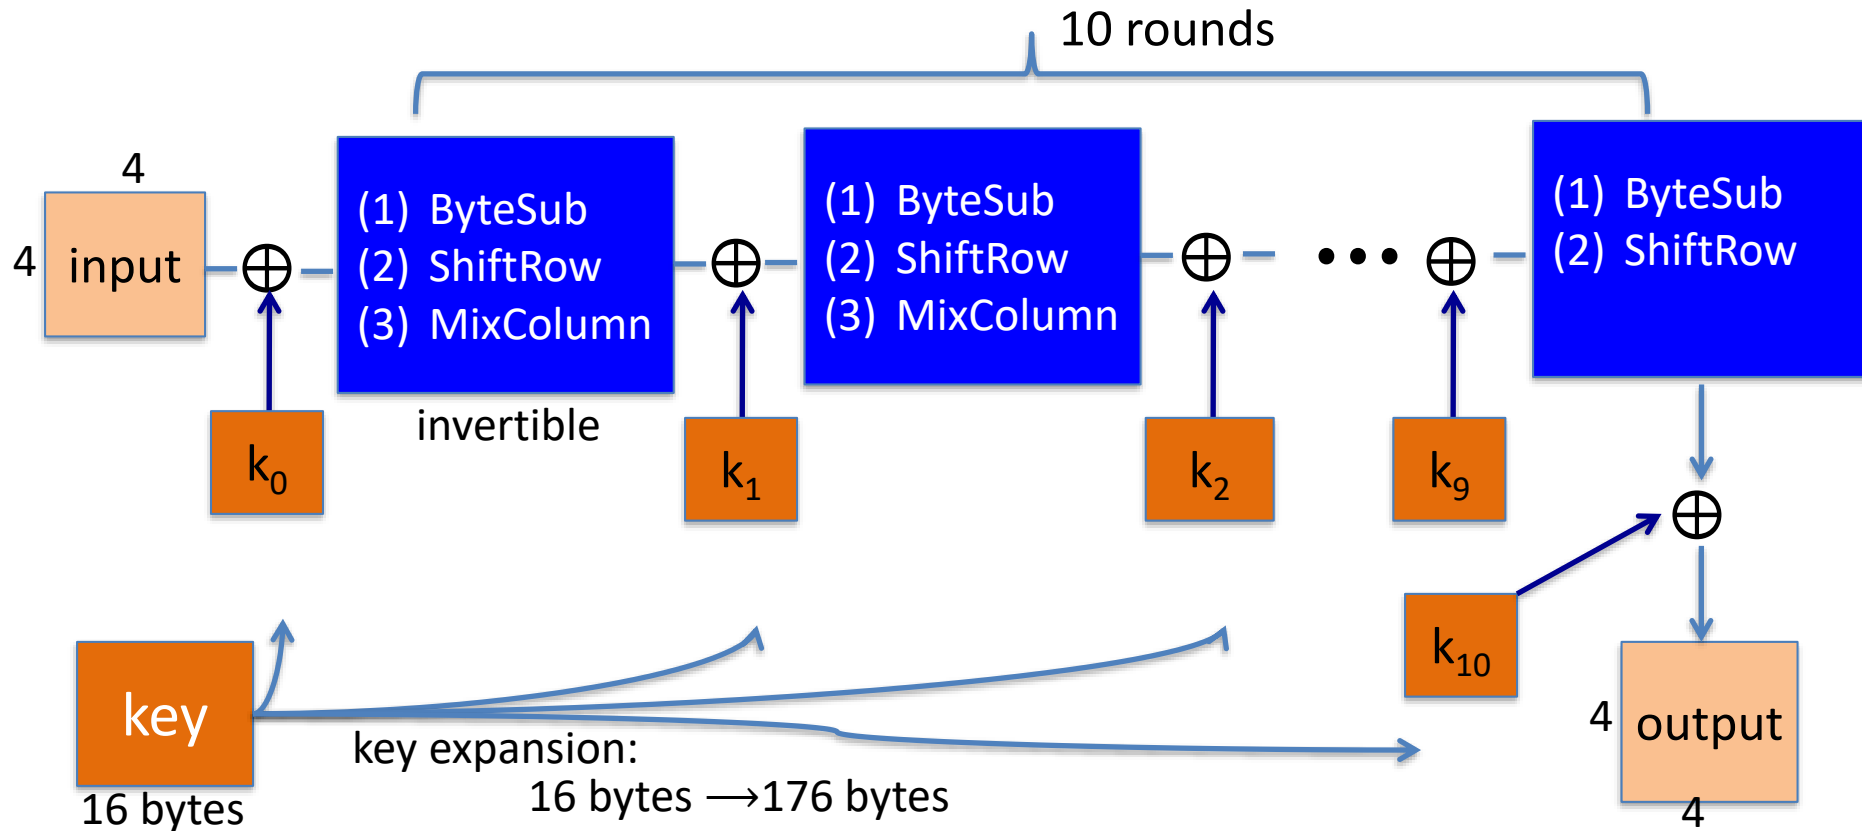
2000: NIST chooses Rijndael as AES (designed in Belgium)

Key sizes: 128, 192, 256 bits. Block size: 128 bits

AES is a Subs-Perm Network (not Feistel)




AES-128 Schematic




The Round Function

ByteSub: a 1 byte S-box. 256 byte table (easily computable)

ShiftRows:



MixColumns:



Code size/performance tradeoff


	Code size	Performance
Pre-compute round functions (24KB or 4KB)	largest	fastest: table lookups and xors
Pre-compute S-box only (256 bytes)	smaller	slower
No pre-computation	smallest	slowest

Example: Javascript AES

AES in the browser:



AES library (6.4KB)
no pre-computed tables



Prior to encryption:
pre-compute tables

Then encrypt using tables

Attacks

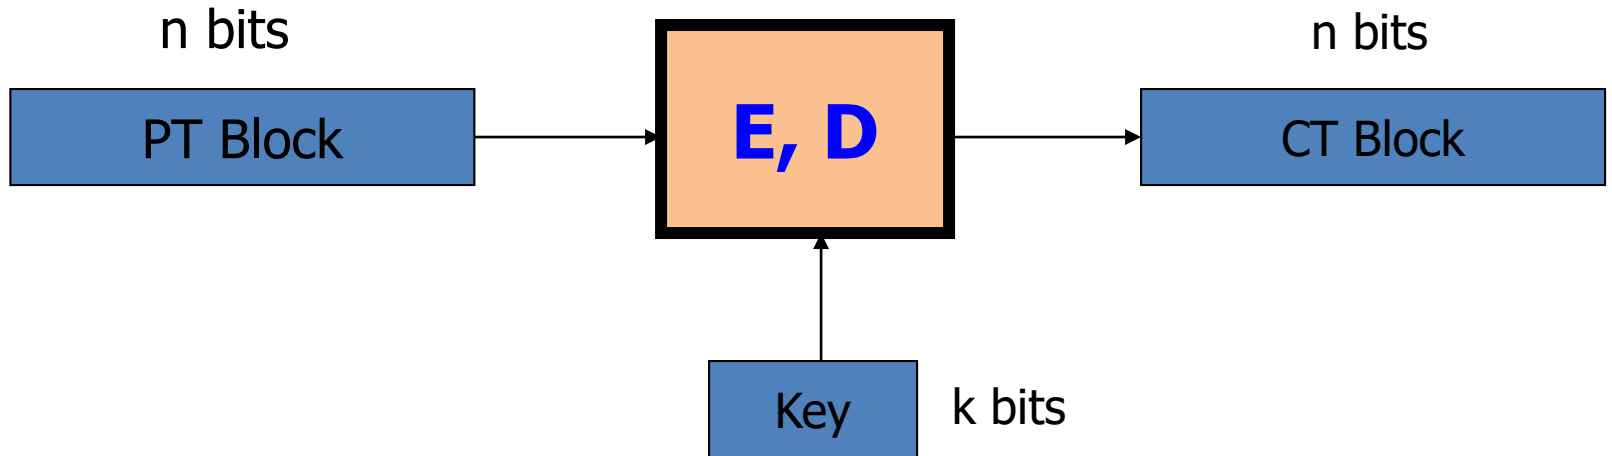
Best key recovery attack:

four times better than ex. search [BKR'11]

Related key attack on AES-256: [BK'09]

Given 2^{99} inp/out pairs from **four related keys** in AES-256
can recover keys in time $\approx 2^{99}$

Review PRFs and PRPs



Canonical examples:

1. 3DES: $n = 64$ bits, $k = 168$ bits
2. AES: $n = 128$ bits, $k = 128, 192, 256$ bits

Abstractly: PRPs and PRFs

- Pseudo Random Function (**PRF**) defined over (K, X, Y) :

$$F: K \times X \rightarrow Y$$

such that exists “efficient” algorithm to evaluate $F(k, x)$

-
- Pseudo Random Permutation (**PRP**) defined over (K, X) :

$$E: K \times X \rightarrow X$$

such that:

1. Exists “efficient” deterministic algorithm to evaluate $E(k, x)$
2. The function $E(k, \cdot)$ is one-to-one
3. Exists “efficient” inversion algorithm $D(k, x)$


Secure PRFs

Let $F: K \times X \rightarrow Y$ be a PRF

$$\left\{ \begin{array}{l} \text{Funs}[X,Y]: \text{ the set of all functions from } X \text{ to } Y \\ S_F = \{ F(k, \cdot) \text{ s.t. } k \in K \} \subseteq \text{Funs}[X,Y] \end{array} \right.$$

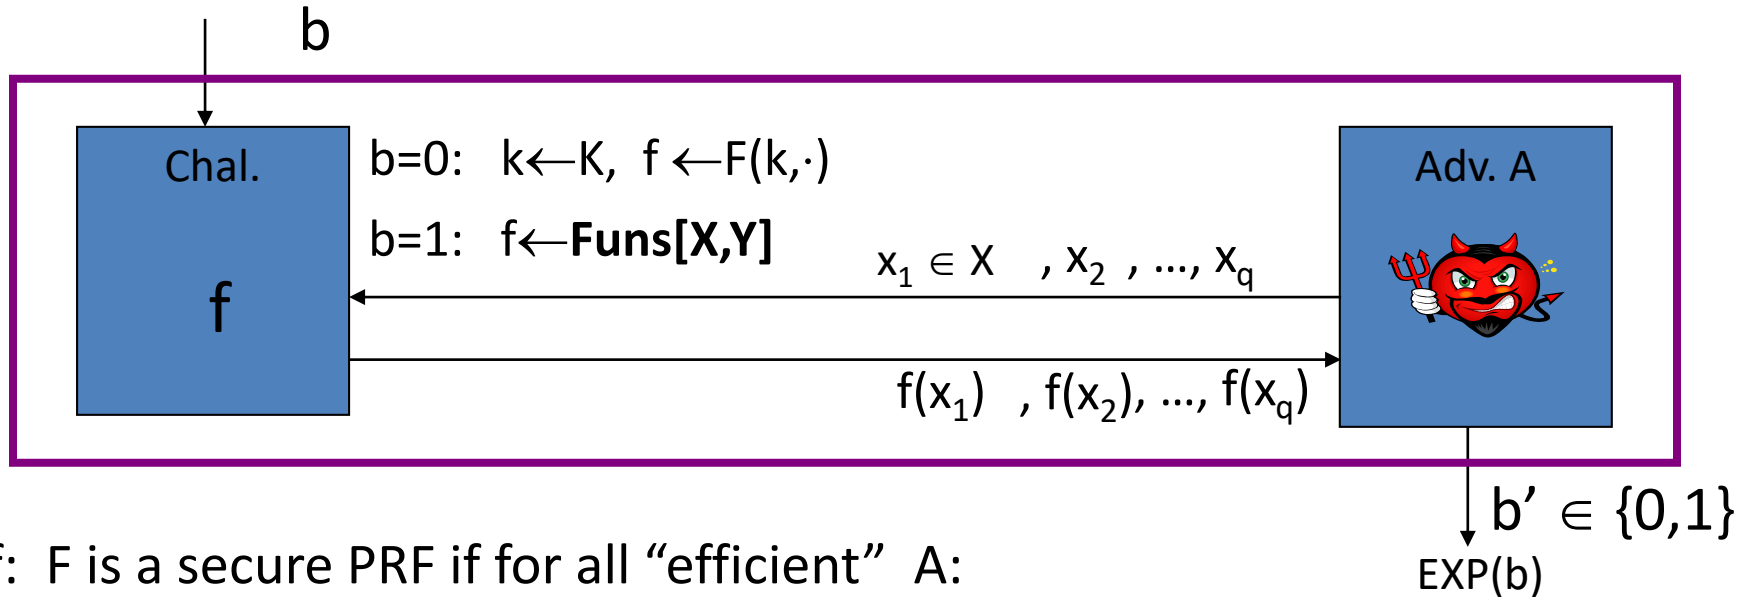
Intuition: a PRF is **secure** if

a random function in $\text{Funs}[X,Y]$ is indistinguishable from
a random function in S_F



Secure PRF: definition

For $b=0,1$ define experiment $\text{EXP}(b)$ as:



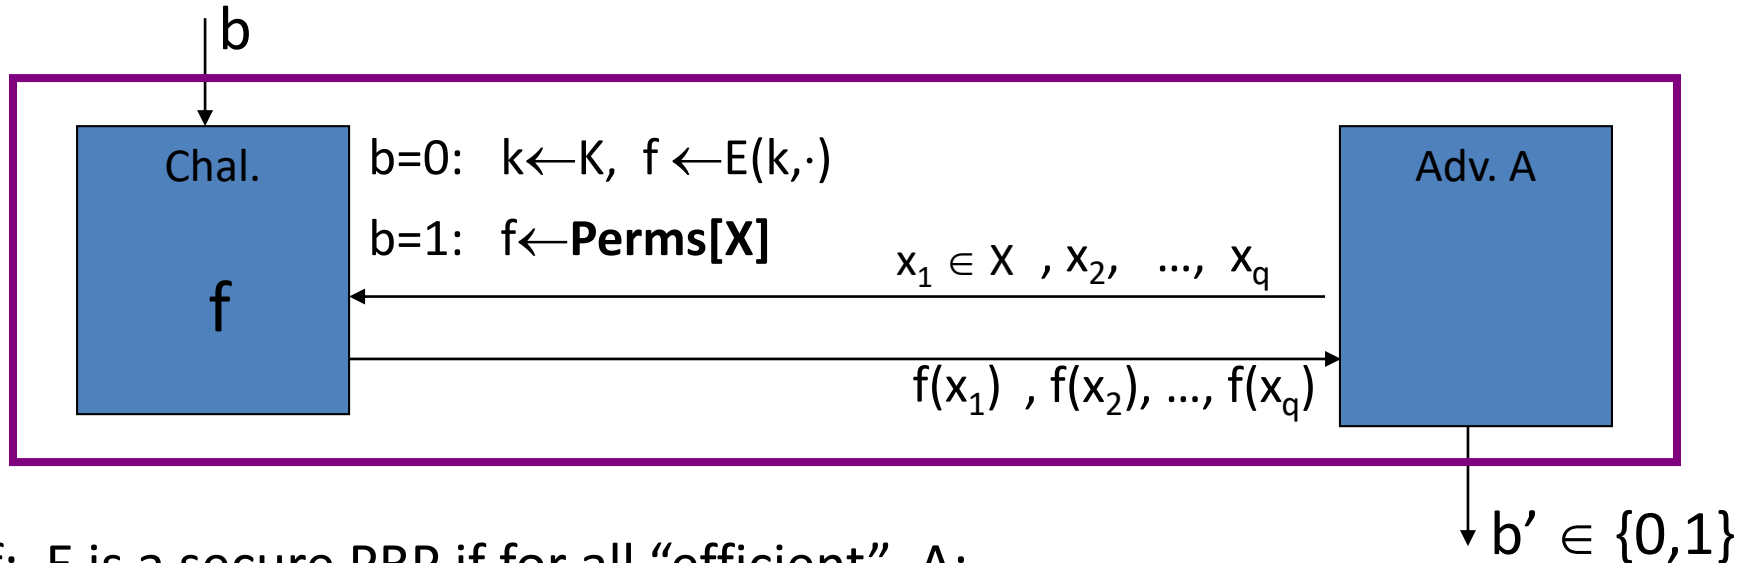
Def: F is a secure PRF if for all “efficient” A :

$$\text{Adv}_{\text{PRF}}[A, F] := \left| \Pr[\text{EXP}(0)=1] - \Pr[\text{EXP}(1)=1] \right|$$

is “negligible.”

Secure PRPs (secure block cipher)

For $b=0,1$ define experiment $\text{EXP}(b)$ as:



Def: E is a secure PRP if for all “efficient” A :

$$\text{Adv}_{\text{PRP}}[A, E] = \left| \Pr[\text{EXP}(0)=1] - \Pr[\text{EXP}(1)=1] \right|$$

is “negligible.”

Example Secure PRPs

PRPs believed to be secure: 3DES, AES, ...

AES-128: $K \times X \rightarrow X$ where $K = X = \{0,1\}^{128}$

An example concrete assumption about AES:

All 2^{80} -time algs. A have $\text{Adv}_{\text{PRP}}[A, \mathbf{AES}] < 2^{-40}$

Final Note

Suggestion:

- don' t think about the inner-workings of AES and 3DES.

We assume both are secure PRPs and will
see how to use them

Security for Many-Time Key

Example applications:

1. File systems: Same AES key used to encrypt many files.
2. IPsec: Same AES key used to encrypt many packets.

Key used more than once \Rightarrow adv. sees many CTs with same key

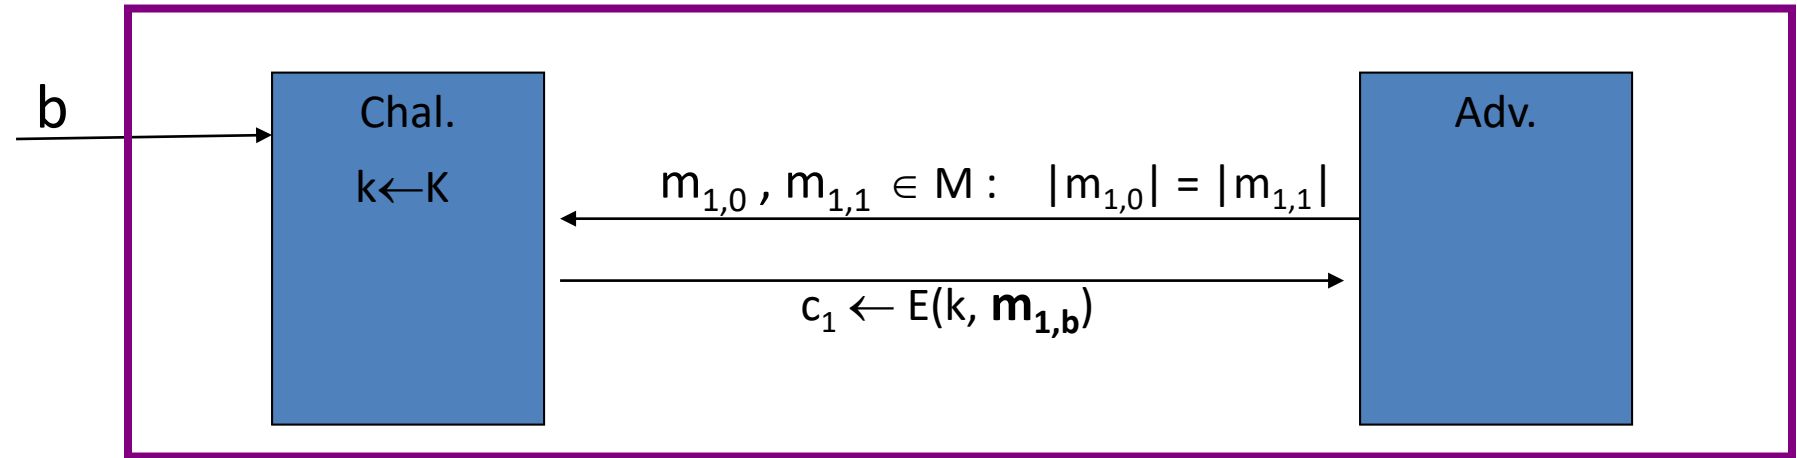
Adversary's power: chosen-plaintext attack (CPA)

Can obtain the encryption of arbitrary messages of his choice
(conservative modeling of real life)

Adversary's goal: Break semantic security


Semantic Security for Many-Time Key

$\mathbb{E} = (E, D)$ a cipher defined over (K, M, C) . For $b=0,1$ define $\text{EXP}(b)$ as:



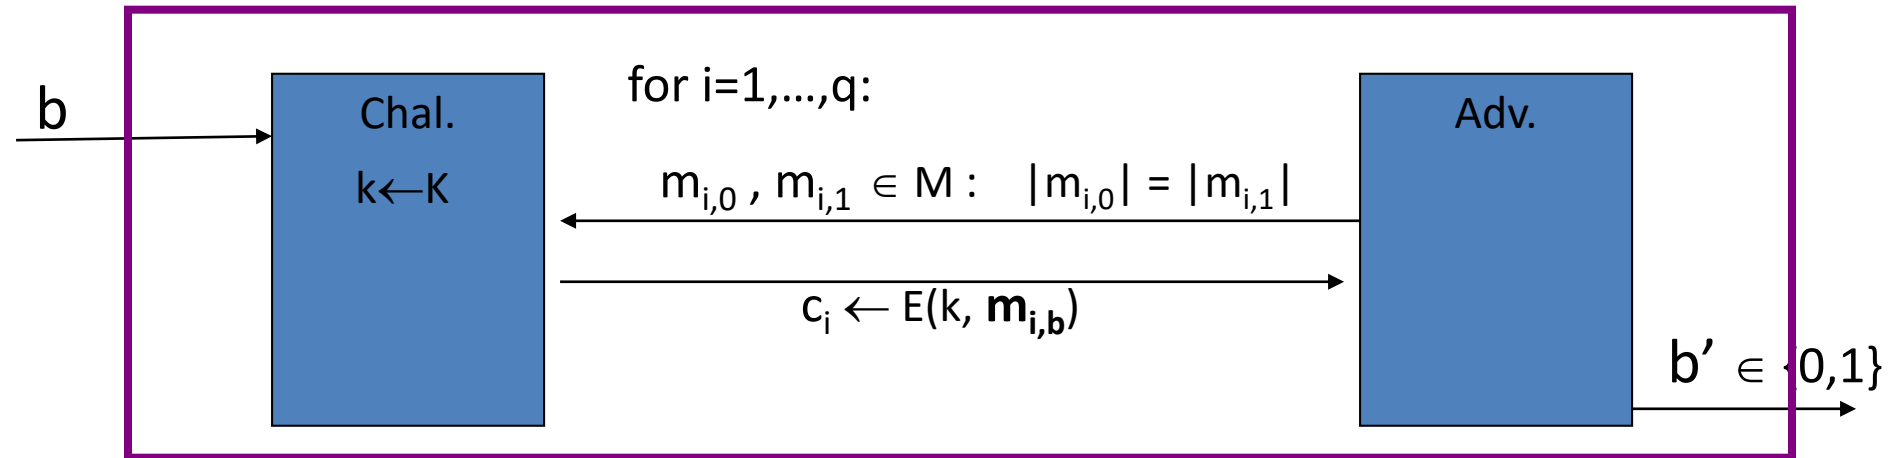
Semantic Security for Many-Time Key

$\mathbb{E} = (E, D)$ a cipher defined over (K, M, C) . For $b=0,1$ define $\text{EXP}(b)$ as:



Semantic Security for Many-Time Key (CPA security)

$\mathbb{E} = (E, D)$ a cipher defined over (K, M, C) . For $b=0,1$ define $\text{EXP}(b)$ as:




if adv. wants $c = E(k, m)$ it queries with $m_{j,0} = m_{j,1} = m$

Def: \mathbb{E} is sem. sec. under CPA if for all "efficient" A :

$$\text{Adv}_{\text{CPA}}[A, \mathbb{E}] = \left| \Pr[\text{EXP}(0)=1] - \Pr[\text{EXP}(1)=1] \right| \text{ is "negligible."}$$

Ciphers insecure under CPA

Suppose $E(k,m)$ always outputs same ciphertext for msg m . Then:




So what? an attacker can learn that two encrypted files are the same, two encrypted packets are the same, etc.

Leads to significant attacks when message space M is small

Ciphers insecure under CPA


Suppose $E(k,m)$ always outputs same ciphertext for msg m . Then:



If secret key is to be used multiple times \Rightarrow
given the same plaintext message twice,
encryption must produce different outputs.

Solution 1: Randomized Encryption

$E(k,m)$ is a randomized algorithm:



⇒ encrypting same msg twice gives different ciphertexts (w.h.p)

⇒ ciphertext must be longer than plaintext

Roughly speaking: CT-size = PT-size + “# random bits”

Solution 1: Randomized Encryption

Let $F: K \times R \rightarrow M$ be a secure PRF.

For $m \in M$ define $E(k, m) = [r \xleftarrow{R} R, \text{ output } (r, F(k, r) \oplus m)]$

Is E semantically secure under CPA?


Yes, whenever F is a secure PRF

No, there is always a CPA attack on this system

→ Yes, but only if R is large enough so r never repeats (w.h.p)

It depends on what F is used

Solution 2: nonce-based Encryption



nonce n : a value that changes from msg to msg.

(k, n) pair never used more than once

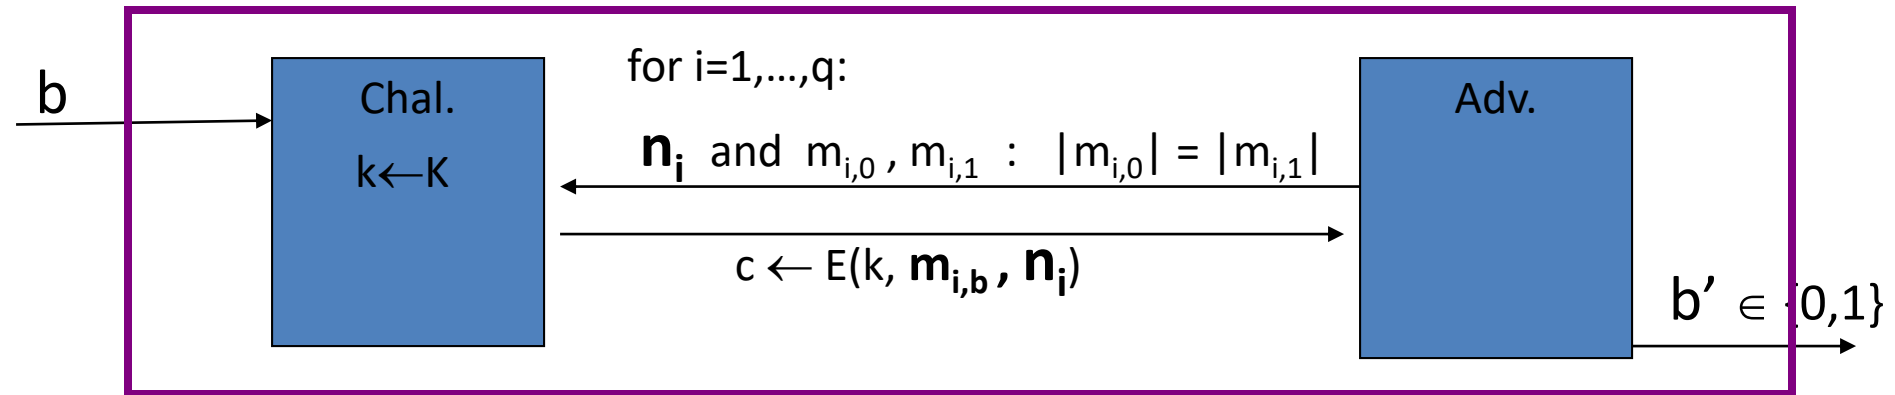
method 1: nonce is a **counter** (e.g. packet counter)

- used when encryptor keeps state from msg to msg
- if decryptor has same state, need not send nonce with CT

method 2: encryptor chooses a **random nonce**, $n \leftarrow \mathcal{N}$

Solution 2: nonce-based Encryption

System should be secure when nonces are chosen adversarially.



All nonces $\{n_1, \dots, n_q\}$ must be distinct.

Def: nonce-based \mathbb{E} is sem. sec. under CPA if for all “efficient” A :

$$\text{Adv}_{\text{nCPA}}[A, \mathbb{E}] = \left| \Pr[\text{EXP}(0)=1] - \Pr[\text{EXP}(1)=1] \right| \text{ is “negligible.”}$$

Lecture 4.4: Message Integrity


Message Integrity

Goal: **integrity**, no confidentiality.

Examples:

- Protecting public binaries on disk.
- Protecting banner ads on web pages.

Message Integrity: MACs



Generate tag:

$$\text{tag} \leftarrow S(k, m)$$

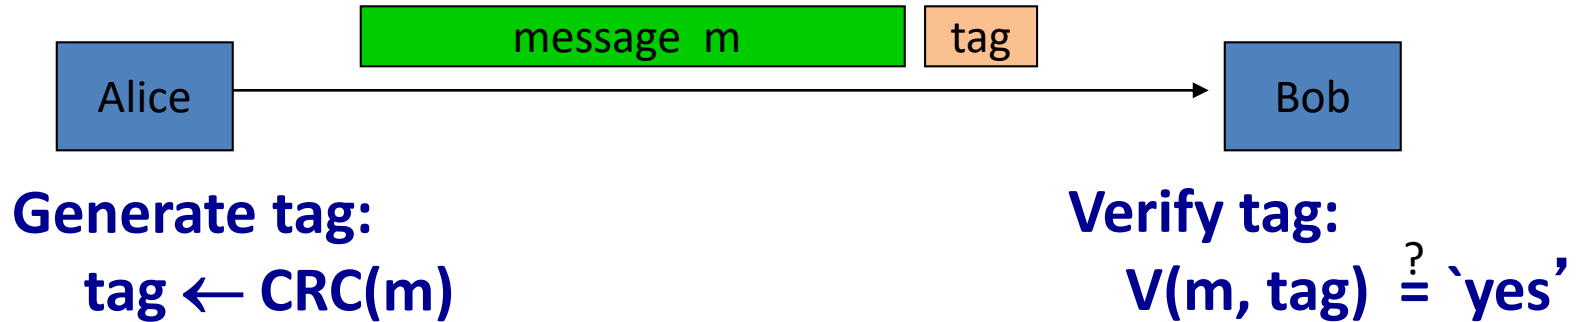
Verify tag:

$$V(k, m, \text{tag}) \stackrel{?}{=} \text{'yes'}$$

Def: **MAC** $I = (S, V)$ defined over (K, M, T) is a pair of algs:

- $S(k, m)$ outputs t in T
- $V(k, m, t)$ outputs 'yes' or 'no'

Integrity requires a secret key



Attacker can easily modify message m and re-compute CRC.

CRC designed to detect random, not malicious errors.

Secure MACs

Attacker's power: **chosen message attack**

for m_1, m_2, \dots, m_q attacker is given $t_i \leftarrow S(k, m_i)$

Attacker's goal: **existential forgery**

produce some new valid message/tag pair (m, t) .


$$(m, t) \notin \{ (m_1, t_1), \dots, (m_q, t_q) \}$$

\Rightarrow attacker cannot produce a valid tag for a new message

\Rightarrow given (m, t) attacker cannot even produce (m, t') for $t' \neq t$

Secure MACs

For a MAC $I=(S,V)$ and adv. A define a MAC game as:



Def: $I=(S,V)$ is a **secure MAC** if for all "efficient" A :

$$\text{Adv}_{\text{MAC}}[A, I] = \Pr[\text{Chal. outputs 1}] \text{ is "negligible."}$$

Example

Let $I = (S, V)$ be a MAC.

Suppose an attacker is able to find $m_0 \neq m_1$ such that

$$S(k, m_0) = S(k, m_1) \quad \text{for } \frac{1}{2} \text{ of the keys } k \text{ in } K$$

Can this MAC be secure?

Yes, the attacker cannot generate a valid tag for m_0 or m_1

→ No, this MAC can be broken using a chosen msg attack
It depends on the details of the MAC

$$\text{Adv}_{\text{MAC}}[A, I] = 1/2$$

Example

Let $I = (S, V)$ be a MAC.

Suppose $S(k, m)$ is always 5 bits long


Can this MAC be secure?

- No, an attacker can simply guess the tag for messages
It depends on the details of the MAC
- Yes, the attacker cannot generate a valid tag for any message

$$\text{Adv}_{\text{MAC}}[A, I] = 1/32$$

Example: protecting system files

Suppose at install time the system computes:



Later a virus infects system and modifies system files

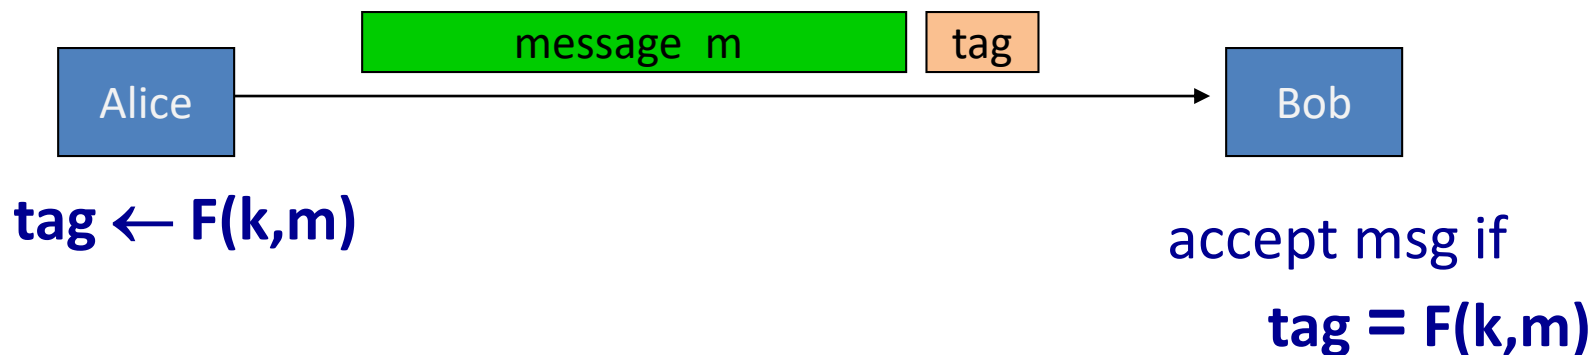
User reboots into clean OS and supplies his password

- Then: secure MAC \Rightarrow all modified files will be detected

Secure PRF => Secure MAC

For a PRF $F: K \times X \rightarrow Y$ define a MAC $I_F = (S, V)$ as:

- $S(k, m) := F(k, m)$
- $V(k, m, t)$: output 'yes' if $t = F(k, m)$ and 'no' otherwise.



A Bad Example

Suppose $F: K \times X \rightarrow Y$ is a secure PRF with $Y = \{0,1\}^{10}$

Is the derived MAC I_F a secure MAC system?

Yes, the MAC is secure because the PRF is secure

→ No tags are too short: anyone can guess the tag for any msg
It depends on the function F

$$\text{Adv}_{\text{MAC}}[A, I] = 1/1024$$

Security

Thm: If $F: K \times X \rightarrow Y$ is a secure PRF and $1/|Y|$ is negligible (i.e. $|Y|$ is large) then I_F is a secure MAC.

In particular, for every eff. MAC adversary A attacking I_F there exists an eff. PRF adversary B attacking F s.t.:


$$\text{Adv}_{\text{MAC}}[A, I_F] \leq \text{Adv}_{\text{PRF}}[B, F] + 1/|Y|$$

$\Rightarrow I_F$ is secure as long as $|Y|$ is large, say $|Y| = 2^{80}$.

Proof Sketch

Suppose $f: X \rightarrow Y$ is a truly random function

Then MAC adversary A must win the following game:



A wins if $t = f(m)$ and $m \notin \{m_1, \dots, m_q\}$

$\Rightarrow \Pr[A \text{ wins}] = 1/|Y|$ same must hold for $F(k,x)$

Examples

AES: a MAC for 16-byte messages.

Main question: how to convert Small-MAC into a Big-MAC ?

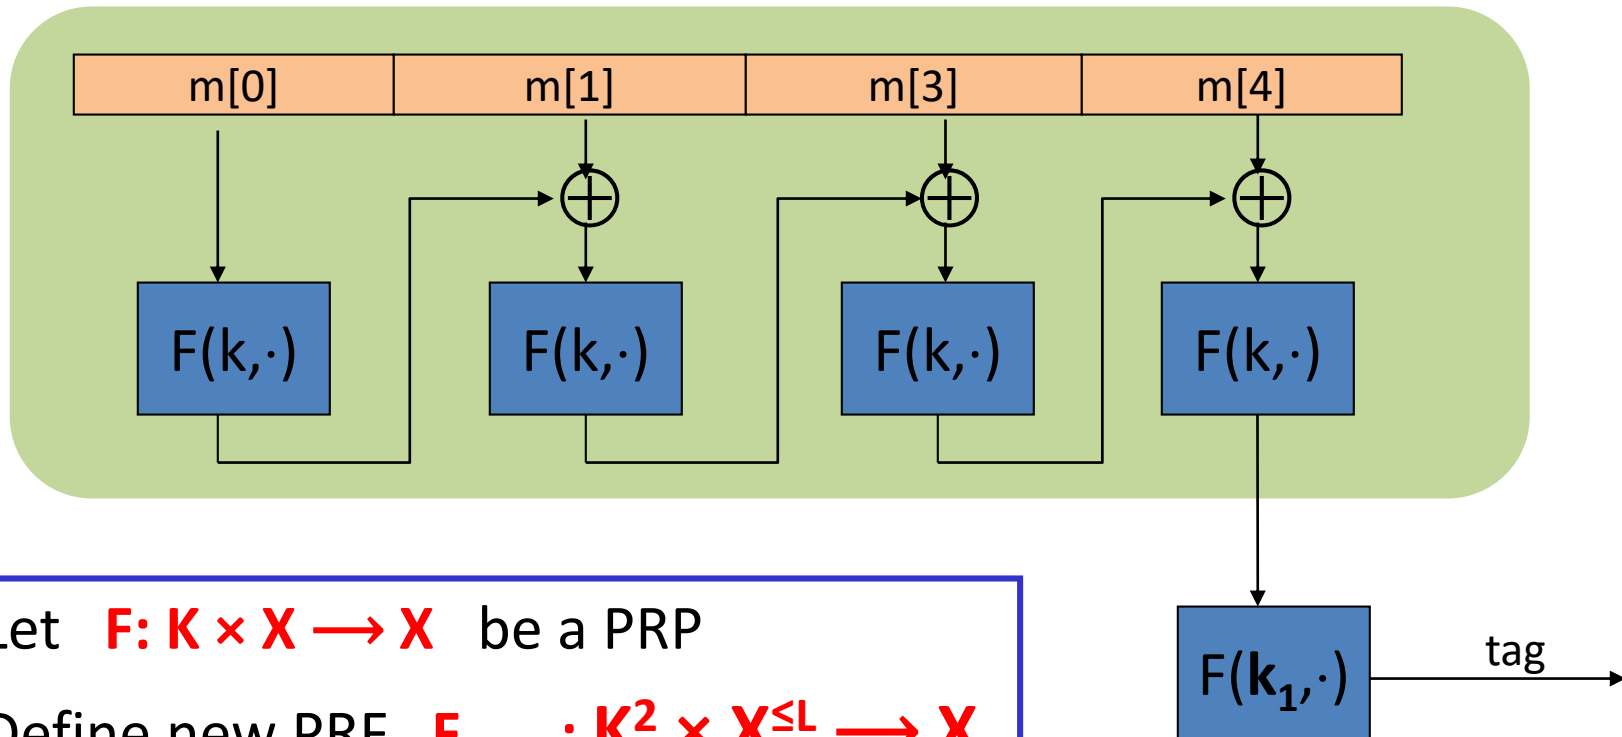
Two main constructions used in practice:

- **CBC-MAC** (banking – ANSI X9.9, X9.19, FIPS 186-3)
- **HMAC** (Internet protocols: SSL, IPsec, SSH, ...)

Both convert a small-PRF into a big-PRF.

Construction 1: encrypted CBC-MAC

raw CBC



Let $\mathbf{F}: \mathbf{K} \times \mathbf{X} \rightarrow \mathbf{X}$ be a PRP

Define new PRF $\mathbf{F}_{\text{ECBC}}: \mathbf{K}^2 \times \mathbf{X}^{\leq L} \rightarrow \mathbf{X}$

Why the last encryption step in ECBC-MAC?

Suppose we define a MAC $I_{\text{RAW}} = (S, V)$ where

$$S(k, m) = \text{rawCBC}(k, m)$$

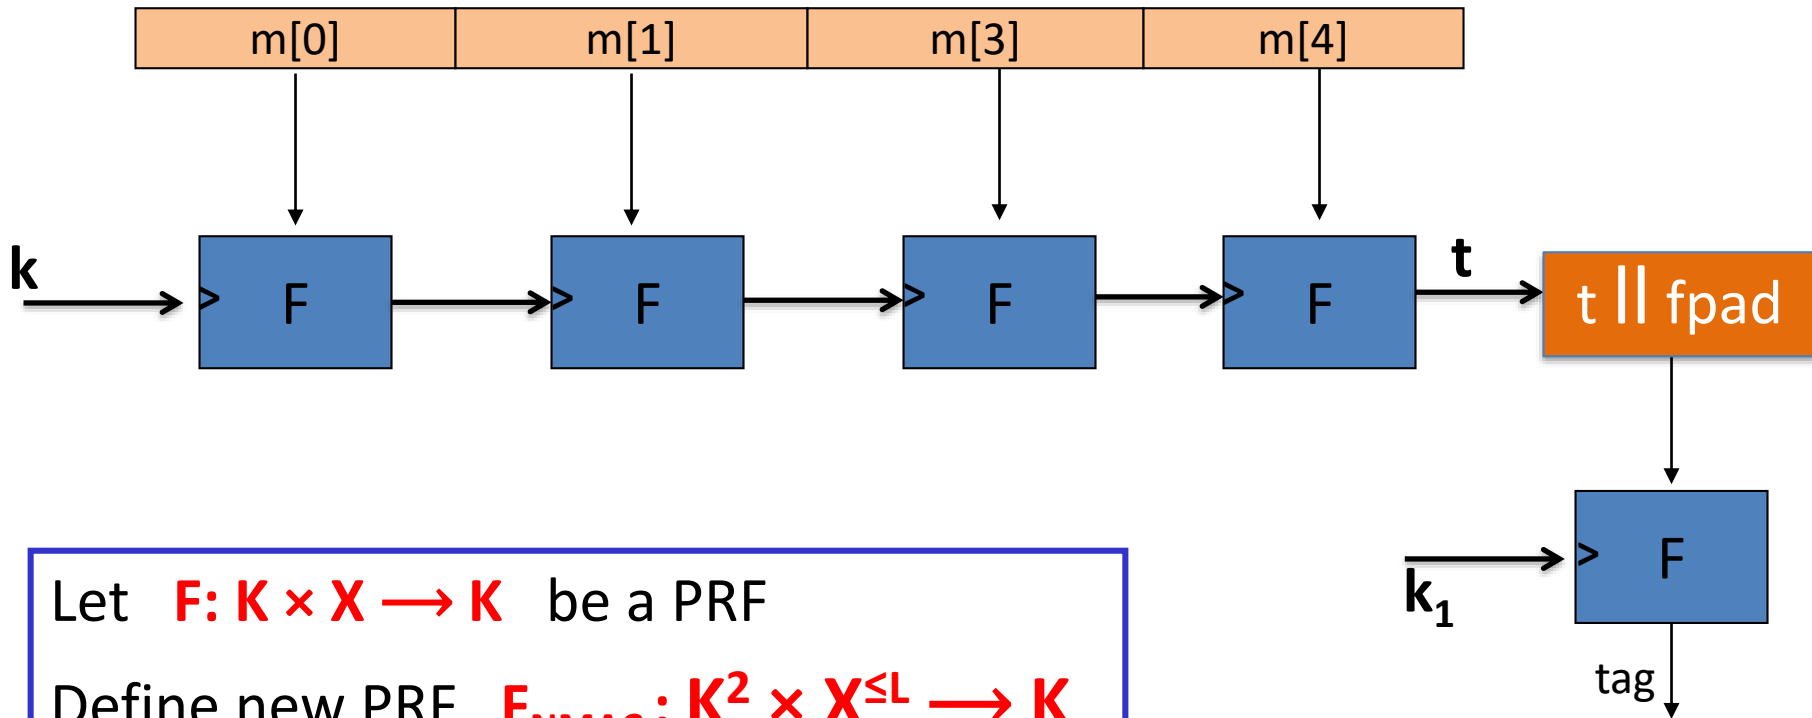
Then I_{RAW} is easily broken using a 1-chosen msg attack.

Adversary works as follows:

- Choose an arbitrary one-block message $m \in X$
- Request tag for m . Get $t = F(k, m)$
- Output t as MAC forgery for the 2-block message $(m, t \oplus m)$

Indeed: $\text{rawCBC}(k, (m, t \oplus m)) = F(k, F(k, m) \oplus (t \oplus m)) = F(k, t \oplus (t \oplus m)) = t$

Construction 2: NMAC (nested MAC)



Truncating MACs based on PRFs

Easy lemma: suppose $F: K \times X \rightarrow \{0,1\}^n$ is a secure PRF.

Then so is $F_t(k,m) = F(k,m)[1...t]$ for all $1 \leq t \leq n$

\Rightarrow if (S,V) is a MAC is based on a secure PRF outputting n -bit tags
the truncated MAC outputting w bits is secure
... as long as $1/2^w$ is still negligible (say $w \geq 64$)

Collision Resistance

Let $H: M \rightarrow T$ be a hash function ($|M| \gg |T|$)

A **collision** for H is a pair $m_0, m_1 \in M$ such that:

$$H(m_0) = H(m_1) \quad \text{and} \quad m_0 \neq m_1$$

A function H is **collision resistant** if for all (explicit) “eff” algs. A :

$$\text{Adv}_{\text{CR}}[A, H] = \Pr[A \text{ outputs collision for } H]$$

is “neg” .

Example: SHA-256 (outputs 256 bits)

MACs from Collision Resistance

Let $I = (S,V)$ be a MAC for short messages over (K,M,T) (e.g. AES)

Let $H: M^{\text{big}} \rightarrow M$

Def: $I^{\text{big}} = (S^{\text{big}}, V^{\text{big}})$ over (K, M^{big}, T) as:

$$S^{\text{big}}(k,m) = S(k,H(m)) \quad ; \quad V^{\text{big}}(k,m,t) = V(k,H(m),t)$$

Thm: If I is a secure MAC and H is collision resistant
then I^{big} is a secure MAC.

Example: $S(k,m) = \text{AES}_{\text{2-block-cbc}}(k, \text{SHA-256}(m))$ is a secure MAC.

MACs from Collision Resistance

$$S^{\text{big}}(k, m) = S(k, H(m)) \quad ; \quad V^{\text{big}}(k, m, t) = V(k, H(m), t)$$

Collision resistance is necessary for security:

Suppose adversary can find $m_0 \neq m_1$ s.t. $H(m_0) = H(m_1)$.

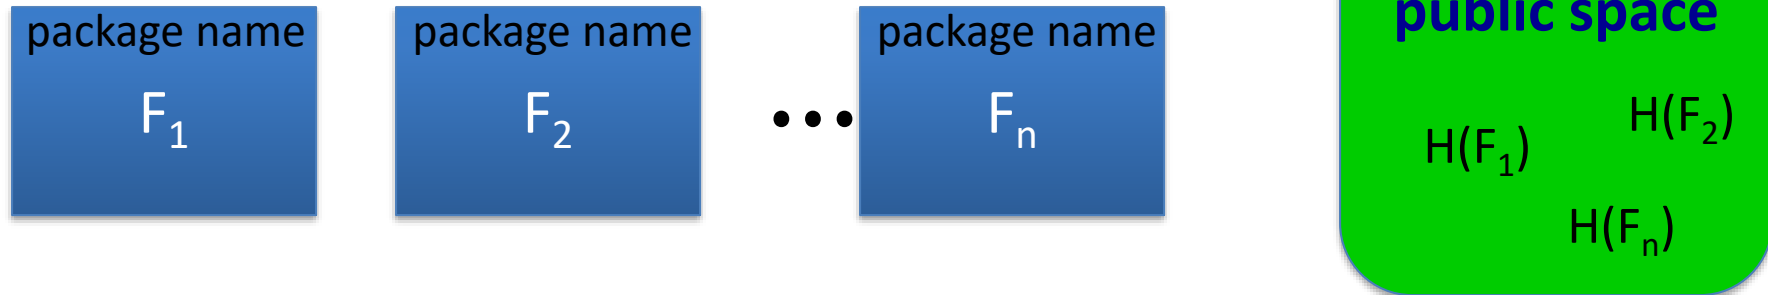
Then: S^{big} is insecure under a 1-chosen msg attack

step 1: adversary asks for $t \leftarrow S(k, m_0)$

step 2: output (m_1, t) as forgery

Protecting file integrity using Collision Resistance Hash

Software packages:



When user downloads package, can verify that contents are valid

H collision resistant \Rightarrow


attacker cannot modify package without detection

no key needed (public verifiability), but requires read-only space

Lecture 4.5: Basic Key Exchange

Key Management


Problem: n users. Storing mutual secret keys is difficult



Total: $O(n^2)$ keys per user

A Better Solution

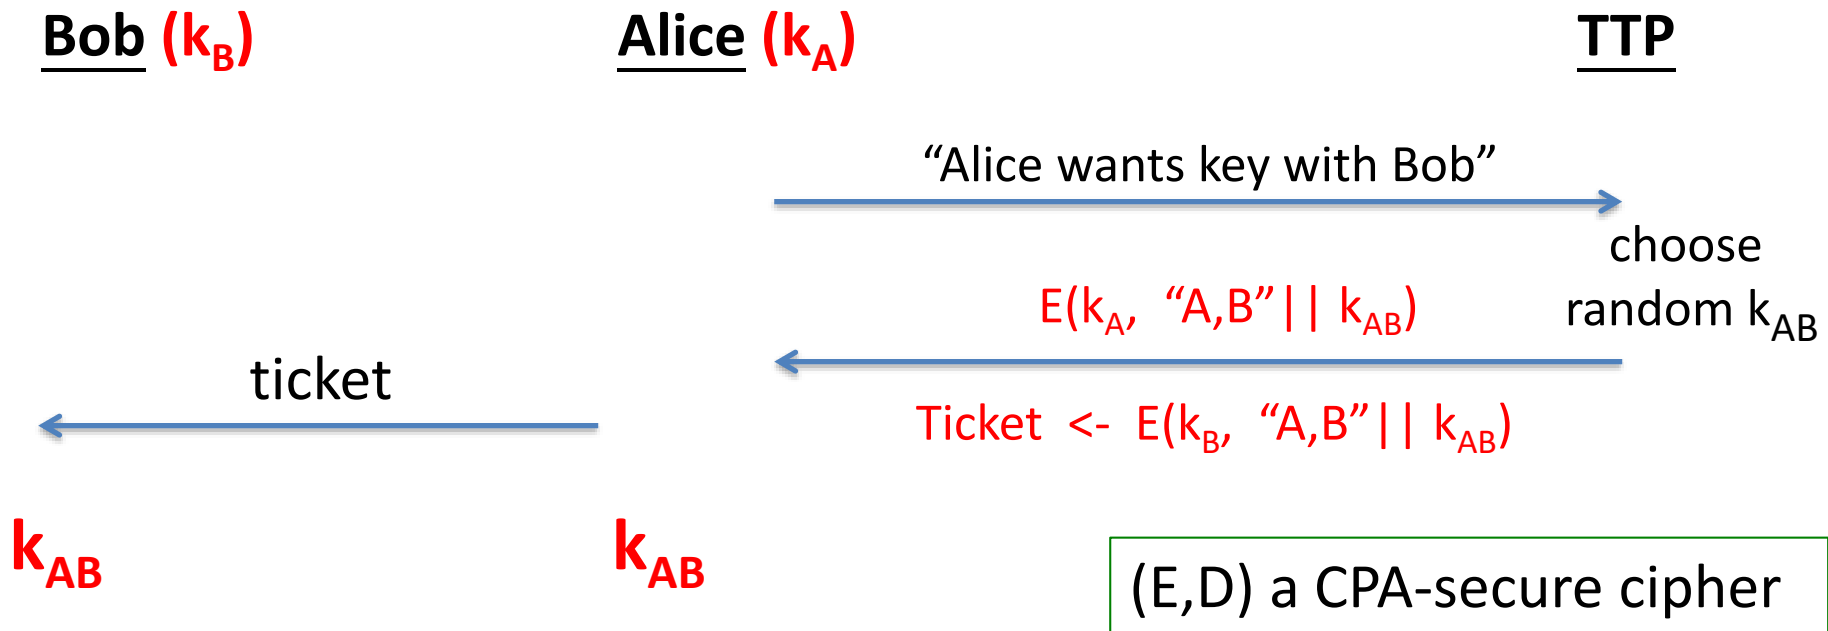
Online Trusted 3rd Party (TTP)



Every user remembers one key

Generating Keys: A toy protocol

Alice wants a shared key with Bob. Eavesdropping security only.



Generating Keys: A Toy Protocol

Alice wants a shared key with Bob. Eavesdropping security only.

Eavesdropper sees: $E(k_A, \text{"A, B"} \parallel k_{AB})$; $E(k_B, \text{"A, B"} \parallel k_{AB})$

(E,D) is CPA-secure \Rightarrow

eavesdropper learns nothing about k_{AB}

Note: TTP needed for every key exchange, knows all session keys.

(basis of Kerberos system)

Key Question

Can we generate shared keys without an **online** trusted 3rd party?

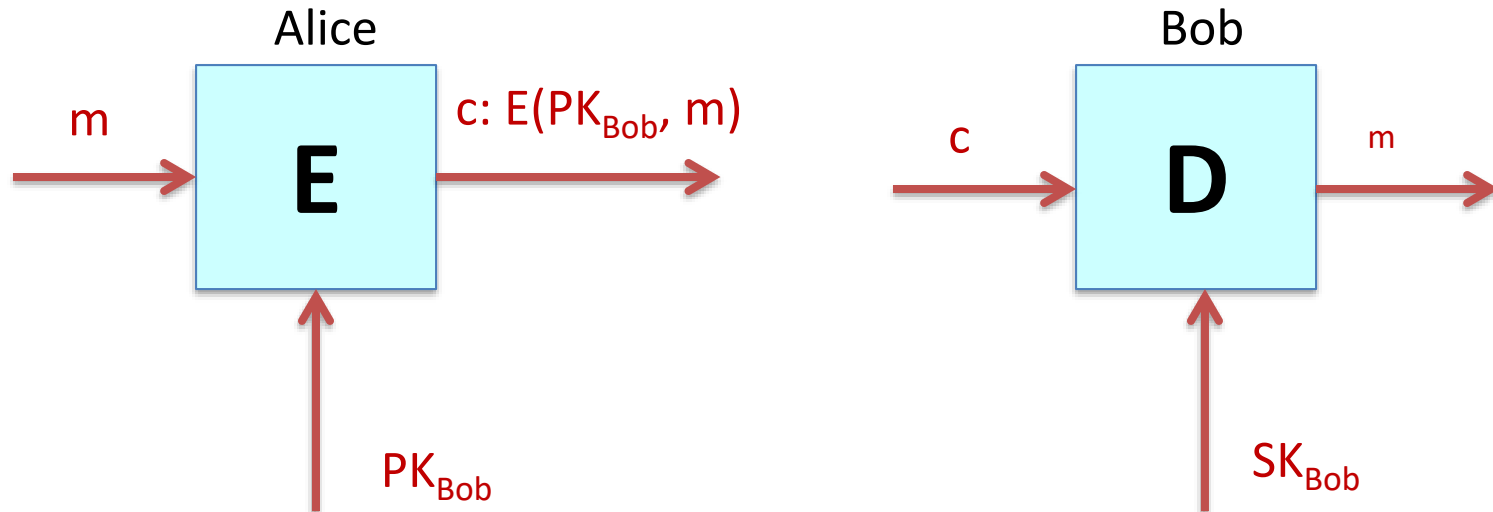
Answer: yes!

Starting point of public-key cryptography:

Merkle (1974), Diffie-Hellman (1976), RSA (1977)

More recently: ID-based enc. (BF 2001), Functional enc. (BSW 2011)

Public Key Encryption



PK: Public Key

SK: Secret Key